OPEN AGENT HACKATHON 2026 · TRACK 3 · REASONING ARCHITECTURE

AgentOffice

An AI office team you just tell what to do. It plans the spreadsheet work, writes the automation, tests it on your data, and asks before anything runs.

Team Automaters · Built on Zetaris, Meterless Nightshift and NVIDIA Nemotron

Office teams lose hours to work nobody has time to fix

WHO

Small office teams

Ops, finance and admin teams with no automation engineer on staff.

WHERE IT HURTS

Repeated manual steps

Email, PDF, spreadsheet, ERP: the same steps many times a day, across tools nobody connected.

WHAT'S MISSING

Someone to fix it safely

Find the routine, prove it's worth it, build and test it, and get sign-off.

SYNTHETIC DATASET · INVOICE MATCHING · HOURS PER WEEK

Three systems, three answers

71.5

Timesheets

What people reported

32.9

Activity log

Measured hands-on time: the one we size automations against

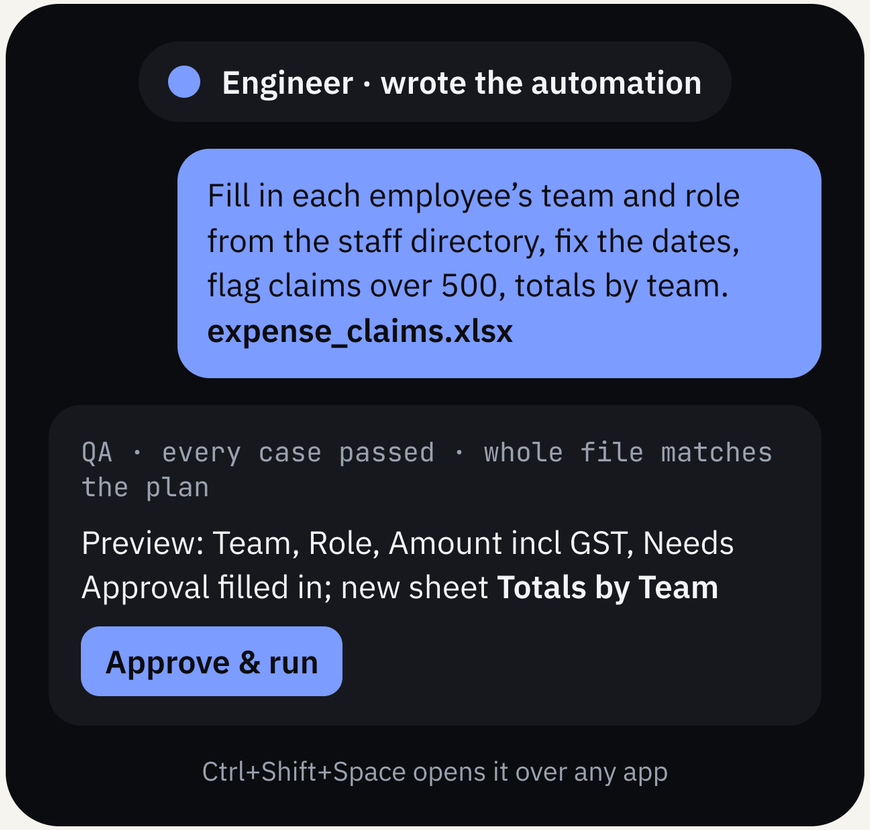
19.5

Ticket system

What people chose to log

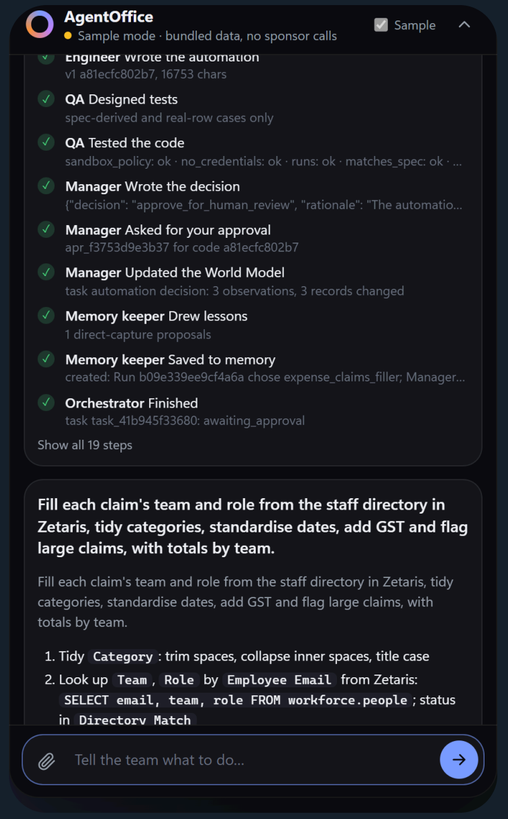
The World Model keeps all three, flags the conflict on the record, and the Conflict Critic explains which source to trust for what.

Tell it the task. Approve the code. Get the file.

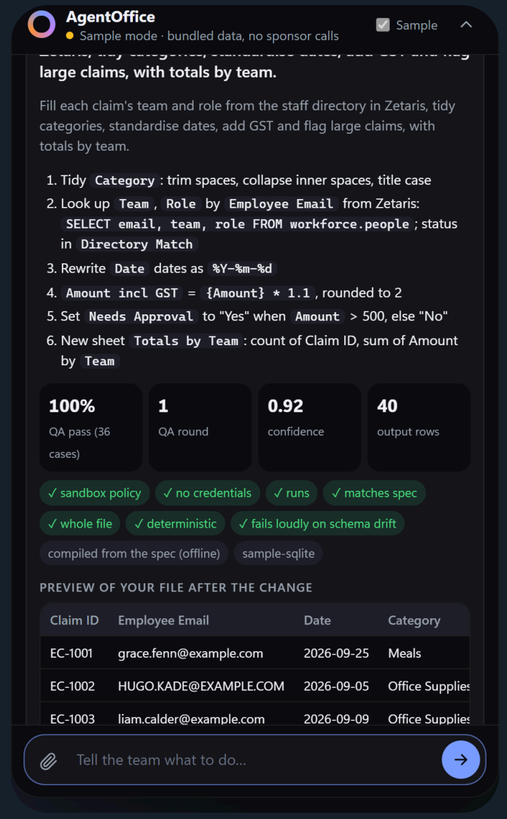


1. **Say it, attach it.** Plain words plus a CSV or Excel file, from a floating island like Copilot.
2. **Plan on real data.** The file is profiled through Zetaris; the PM writes a step plan; a critic checks every column, lookup and date format against it.
3. **Build and test.** Nemotron writes the Python; QA checks it against an independent reference of the plan on your rows and edge cases, and sends failures back.
4. **You approve the exact code.** Only then does it run: a filled workbook with an audit sheet.
5. **It remembers.** Next week’s file: same approved code, re-tested, runs without asking.

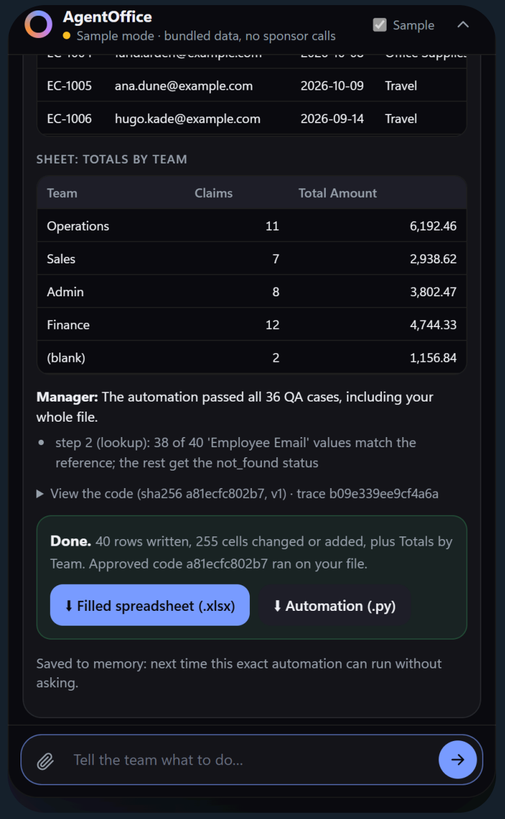
What it looks like



Every agent step streams in live



The plan, QA checks and a preview before approval



Approved code runs; download the filled file

Your file is ready done

Screens from the expense-claims demo in sample mode; live runs show the Nemotron route in the header.

Ask where the time goes: three stages, one shared memory

1 · Discover · Track 1

Where does the time go?

Source Mapper · Reconciler · Conflict Critic

2 · Explain · Track 2

Why is this the bottleneck?

Research Planner · Retriever · Report Writer · Insight Critic

3 · Build · Track 3, primary

What should we do about it?

PM · Priority Critic · Engineer · QA · Manager · Memory Keeper

Shared World Model: every agent reads its context from it and writes what it learns back. **H-MEM** carries lessons and human decisions from one run to the next.

The build loop argues first

Priority Critic

Sends back any claim above the hands-on time the activity log measures · max 2 rounds

PM
Ranks routines, writes the spec

Engineer
Writes code from the spec

QA
Tests the code against the spec

Manager
Decides behind a rule gate

Human
Approves the exact code hash

QA’s failing cases go back to the Engineer · max 3 revisions

Every arrow is a traced handoff in logs/trace-<run_id>.jsonl, hash-chained so edits are detectable.

QA tests the spec, never the code

1. **The PM writes a spec, not prose**

   A small rule language: a key, ordered checks, and the statuses they lead to.

2. **QA derives the expected results**

   From the spec alone: every branch, tolerance edges, case and spacing, ambiguous keys, and a renamed column that must fail loudly.

3. **The code runs in a sandbox**

   Allowed imports only, no eval or file access, an empty environment and a hard timeout.

```
invoice_po_matcher
key: po_number (strip, upper)
check: vendor equals_ci → vendor_mismatch
check: amount within 1% → amount_mismatch
statuses: no_po · ambiguous_po · matched
```

> Halve the tolerance in the code, and QA catches it. That case is in the test suite.

Memory changes what the team does next

Session 1

Learns

The team proposes the invoice matcher. The finance manager rejects it: “ERP writes need my sign-off first.”

Session 2

Recalls

The Memory Keeper retrieves that decision. The PM proposes order-status replies instead.

Session 3

Acts on it

With that approved, it moves on to CRM dedupe. A Finance-only question escalates instead of asking again.

H-MEM, built to the Meterless reference: 8-signal retrieval with reasons, an approval boundary for new memories, conflicts resolved on the record, a hash-chained trust ledger. Tested end to end.

Zetaris is the data layer, end to end

Your upload is a Zetaris table

Our API serves the workforce tables and every uploaded sheet (one row per cell) as JSON; seven statements, run one at a time after a dry run, register them as Lightning REST tables.

Agents write plain SQL

The profiler reads your file with SQL, the PM’s lookups pull the staff directory with SQL, and the analysts ground every finding in SQL rows. No views; the trace keeps both versions of the SQL.

Verified on a live install

All 15 pipeline queries ran on local Zetaris 2.4.3, with findings identical to the sample run.

REST table	Rows
people	12
timesheets	827
tickets	223
email_events	2,672
task_events	8,266
upload_cells	live

Workforce counts match the generator’s manifest; upload_cells is re-read at query time, so a file is queryable the moment it’s uploaded.

Meterless: provider and two engines

Power rule

Nightshift, primary provider

The full provider contract: status from GET /v1/models, auto-detect, connect links, a 180 s timeout. Every call logs the model that answered.

Build rule · engine 1

World Model

Content-hash IDs, an append-only log with row provenance, a 10-step idempotent ingest. Stage 1’s source conflicts are its conflict records.

Build rule · engine 2

H-MEM

Tiered memory behind an approval boundary, 8-signal retrieval with reasons, a hash-chained trust ledger. Approved automations run again without asking.

NVIDIA: Nemotron models run through Nightshift lanes, hosted, with no local GPU. Route used: meterless/deep → nvidia/nemotron-3-ultra-550b-a55b, recorded in every trace.

Every step leaves evidence

Agent	Action	Status
SourceMapperAgent	zetaris_query	skipped (sample)
ConflictCriticAgent	critique_sources	needs_attention
InsightCriticAgent	critique_claims	approved
PriorityCritic	critique_priority	approved
QAAgent	critique_output	approved
ManagerAgent	request_approval	pending

Real events from a sample-mode run:
output_examples/sample_mode/example_1.trace.jsonl (59 events).

Hash-chained

Edit or delete a line and verification fails.

Bound to code

An approval covers one code hash, nothing else.

Runs without keys

Sample mode runs the real agent logic on synthetic data.

158 tests

Also passing inside the Docker image on Python 3.11.

Live results: Nemotron + Zetaris, 25 Oct

Spreadsheet task	Plan steps	QA rounds	Final QA · after approval
Expense claims: staff lookup in Zetaris, dates, GST, flags, totals	7	3 (policy, then blanks)	49 / 49 · 275 cells
Leads vs CRM sheet: split names, owner, status	4	1	34 / 34 · 119 cells
Vendor invoices: dedupe, fill, overdue, sort	6	1	40 / 40 · 34 → 30 rows

Analysis question	Automate first	Saved / wk	QA
Finance: where are we losing time?	invoice matching (high risk)	≈ 24.4 h	28 / 28
Operations: what is eating the week?	order status updates	≈ 7.9 h	25 / 25
Sales: is CRM data entry worth it?	CRM data entry	≈ 7.1 h	28 / 28

What it’s built with

Runtime

Python 3.11, FastAPI, httpx, openpyxl; a plain-JS chat island

State

SQLite for the World Model and H-MEM; JSONL traces

Data

Zetaris 2.4.3 in local Docker, Lightning REST tables

Models

Meterless Nightshift, NVIDIA Nemotron lanes

Packaging

Docker on python:3.11-slim; sample mode needs no keys

Built with

Claude Code; every line reviewed and tested

What comes next

1 **Real exports in**

Register a team’s ERP and CRM exports in Zetaris; the agents’ SQL stays the same.

2 **Dry run, run and undo**

Show the changes an approved automation would make, then apply them reversibly.

3 **Scheduled runs and more step types**

Run approved automations on each new file automatically; add steps for emails, PDFs and pivots.

4 **Memory that consolidates**

H-MEM dreaming and sleep: merge related lessons, archive stale ones.

5 **Priorities learned from approvals**

Weight the PM’s ranking by what humans approved and rejected before.

A team that argues, tests and remembers, and asks before it acts.

github.com/allurimohansrisaivarma-coder/AgentOffice

```
docker run -p 8000:8000 -e SAMPLE_MODE=true oah-submission
```

Team Automaters · Track 3, Reasoning Architecture · Open Agent Hackathon 2026